


Reranking quality against its latency cost



P95 of the per-query reranking call, measured in the run itself; one-off model load time is excluded (it is reported as model_load_seconds in reranker_metrics.json and can reach several minutes). The x axis is logarithmic -- equal distances are equal RATIOS, not equal milliseconds. On the frontier: original_ranking@10, fuse_bge_first@10. Marks with an ink ring are the baseline and the winner.