


Is a deeper candidate pool worth its cost?



Top: the ceiling rises only +0.013 Hit@1 across the whole depth range in candidate_ceiling.csv, so deeper pools add very little headroom; the measured rerankers (coloured) show what is actually realised. Bottom: the same depths on a log latency axis. A depth with no marker for a given reranker was not scored at that depth in these run dirs. Latency excludes one-off model load. Omitted from the latency panel because their P95 is below 0.05 ms at every depth (they would add five empty decades to the log axis): original_ranking.